Artificial Intelligence–Driven X-Ray Analysis: Enhancing Image Quality and Diagnostic Accuracy

1 Chitkara University Institute Of Engineering And Technology, Punjab Email: nandinibakshi4@gmail.com

2 Chitkara University Institute Of Engineering And Technology, Punjab Email: gnitya136@gmail.com

3 Chitkara University Institute Of Engineering And Technology, Punjab Email: khushibansal995@gmail.com

4 Chitkara University Institute Of Engineering And Technology, Punjab Email:

monika.aggarwal@chitkara.edu.in

Abstract:

Through medical imaging, it is possible to find and be diagnosed with diseases, especially for patients who have a respiratory condition like pneumonia. Early diagnosis is extremely important in improving the overall outcomes for a patient. There are many ways of doing medical imaging; however, of the different types of imaging techniques, chest x-rays are the most common imaging technique used because they are low cost, widely available, and easy to complete. One of the biggest problems with chest x-ray images is the difficulty in accurately interpreting those images; due to differences in image quality and contrast, as well as the number of subtle abnormalities, it is often very difficult to tell the difference between a normal and an abnormal image when viewed manually. To solve some of the limitations associated with manually identifying chest x-ray images, we propose using an artificial intelligence (AI)-based imaging system that improves the quality of chest x-ray images, while increasing the diagnostic accuracy of those efforts.For realizing these objectives through this research, some main objectives have been set. Firstly, there is an attempt to improve the quality of the image for the purpose of analyzing the data effectively. Then, there is an objective that seeks to develop a strong classification model with the help of deep learning approaches in order to detect any kind of anomalies. The proposed system integrates advanced techniques for preprocessing images to enhance their quality, along with an AI-based deep learning model using the ResNet-18 architecture, and transfer learning to create an accurate predictive model. The model was developed using a publicly available data set, and then validated using a separate data set, in order to determine if the model had the ability to generalize and would be robust to use in the future. The results from the experiments show that the proposed model is highly accurate and has a good recall rate when distinguishing normal cases from abnormal images, and is a potential method of assisting in the early-diagnosis of patients. Grad-CAM visualizations were also generated during the experiments and were used to assist with interpretability of the AI predictive model's results and decision-making process; which increase transparency and improve trust in the use of the model for clinical applications. This article discusses and highlights.

Keywords:

Artificial Intelligence, Deep Learning, Chest X-ray Analysis, Pneumonia Detection, ResNet-18, Image Enhancement, Transfer Learning, Grad-CAM, Medical Image Classification, Explainable AI.

Introduction:

Chest X-ray images have become an essential component of contemporary medicine, especially in diagnosing respiratory diseases. They consistently rank among the top ten most utilized forms of imaging because they are both low-cost and widely available (through both hospitals and clinics) and easy to take. However, common challenges exist for radiologists when it comes to interpreting chest X-rays, and these include (but are not limited to) the ability to isolate minute pathologies as well as distinguishing things such as stackable anatomy and image quality.

Further, there is often significant variation in how many chest X-rays a radiologist will review on an average day, meaning that fatigue, human error, or delays may create significant disadvantages for patient treatment and care. In light of these issues, artificial intelligence (AI) has emerged as a feasible way for radiologists to aid them in interpreting their workload of medical images. As an example, deep learning models utilized to identify abnormal patterns and improve image quality or offer faster preliminary diagnoses can significantly reduce the burden placed on radiologists. As such, the consistent employment of AI will help radiologists improve the consistency and accuracy of their interpretations. Practical applications for this system within healthcare include assisting radiologists with both fast diagnoses, and the efficient screening of chest related diseases within hospitals. For those living in rural or isolated areas, where they have limited access to specialists such as radiologists, this system would supply initial diagnostic assistance. The system is also very useful in emergency situations when time is of the essence and decisions regarding clinical care need to be made quickly. Additionally, integrating this system with telemedicine services will improve access to healthcare by enabling remote diagnosis and consultation; especially for people who live in under-served areas. With the goal of developing an intelligent machine learning system that can help establish improved chest X-ray images, accurately detect and display abnormalities, and provide clear explanations of its recommendations, this research project is designed to support, not replace healthcare practitioners in making quicker, more accurate, and clinically sound decisions. Enhancing existing medical imaging processes by incorporating AI technologies into these processes will allow for higher accuracy in diagnosis, earlier detection of disease, and more reliable and efficient clinical decision-making, leading to improved patient outcomes. Overall, the findings of this research will help to develop an area of trustworthy explainable AI systems in the area of healthcare.

Literature Review:

The rise in popularity of deep learning over recent years has been considerable, particularly within the realm of medical imaging. Convolutional Neural Networks (CNN's) are capable of learning highly complicated patterns from input images and have exhibited exceptional performance when it comes to subsequent classification tasks. Numerous researchers have investigated the use of CNN based architectures such as VGG, ResNet , and DenseNet for disease detection from X-ray images. One critical aspect of the continued advancement of CNN based models is the concept of transfer learning, whereby an already trained CNN (on a large dataset) can be utilized to perform a different but related function. Transfer learning allows researchers to produce satisfactory results using only small quantities of medical image data.

Another crucial consideration regarding the successful implementation of CNN based models in the medical imaging realm is image preprocessing. Preprocessing techniques such as CLAHE can improve the contrast of images, thus facilitating the detection of previously hidden features. Furthermore, new techniques such as Grad-CAM have been developed to improve the transparency of AI based predictive models, allowing researchers to determine which elements of an input image are most responsible for generating a given prediction.

Although all of the aforementioned techniques have been shown to be effective at some level, several challenges still exist, including but not limited to overfitting, inability to interpret results, and false predictions. Recent studies on the processing of medical imaging data through deep learning models have advanced significantly, particularly with regard to increasing their robustness and ability to generalize by combining preprocessing techniques with deep learning architectures. Emerging hybrid models that employ their effective combination of image enhancement and transfer learning techniques have produced highly encouraging improvements in the performance of classification results.

Methodology:

a. Dataset and Preprocessing

Gaussian Blur applied for noise reduction CLAHE used to enhance contrast [4] We did edge sharpening to bring out important details All images resized to 224x224 and normalized which in turn improved the quality of the input data. We divided the data into training, validation and test sets to properly evaluate the model.

The training set for pattern learning, the validation set for tuning model parameters. And the test set which did not see training was used to see how well the system performed.

b. Model Architecture

The model used in this study is ResNet-18, [1] which is one of the most popular deep learning architectures. Instead of training this model from scratch, transfer learning is used. The final layer of this model is replaced with a new layer that is appropriate for binary classification. This enables the model to classify images into normal and abnormal images.

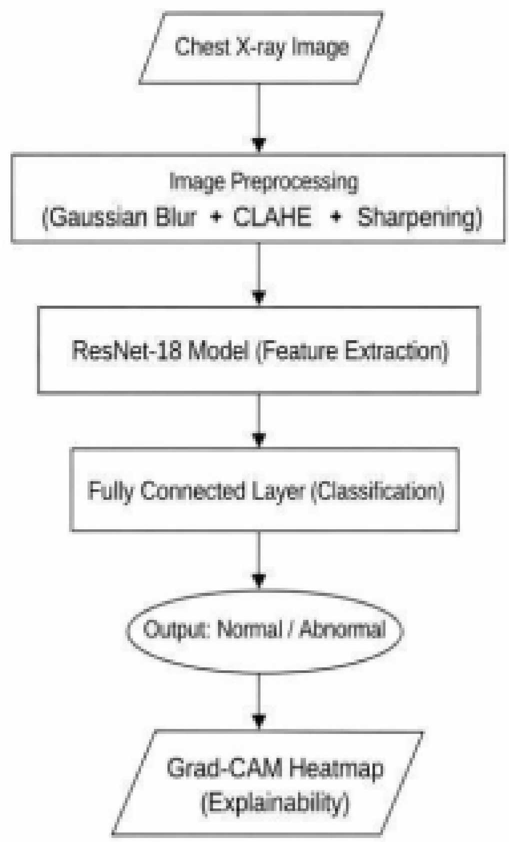

c.Training Process

The training for the model occurs via the PyTorch Framework. To have an effective way to evaluate the model, the dataset was partitioned into three categories; the training dataset, validation dataset, and testing dataset. The selection of key hyperparameters, such as learning rate, batch size, and the number of epochs which were used during training were specifically determined to build robust classification models, using Adam to optimize model weights and utilize binary cross-entropy for the loss function. In order to obtain generalization capabilities and to reduce the negative effects of over-fitting, the training dataset underwent augmentation that included rotation and horizontal flipping of the images. Additionally, preprocessing of the images includes normalization and resizing to ensure consistency with the input data to the model. Early stopping and model checkpointing techniques were also employed to monitor the validation performance of the model during training and retain the best performing version of the model. Throughout the training of the model, backpropagation was utilized to adjust/modify the weight (parameters) of the model to minimize the loss of each epoch of training. At the conclusion of multiple epochs for training, there was an observable improvement in the accuracy and loss of the model network, as demonstrated in figures 1 and 2. In conclusion, the training process was designed to create a model which will provide high accuracy and reliability for the classification of chest X.composite images

(X-ray’s).

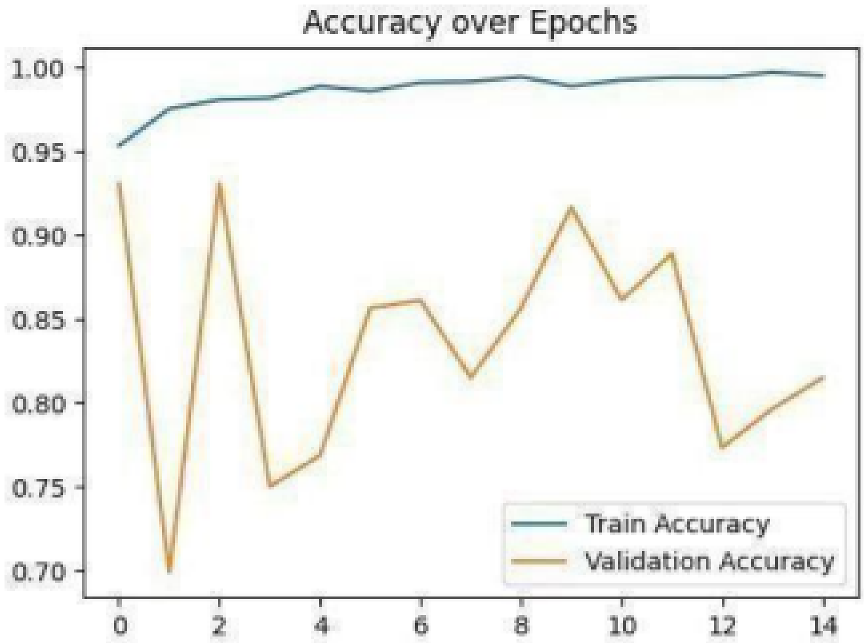

Fig. 1. Training and validation accuracy of the proposed model

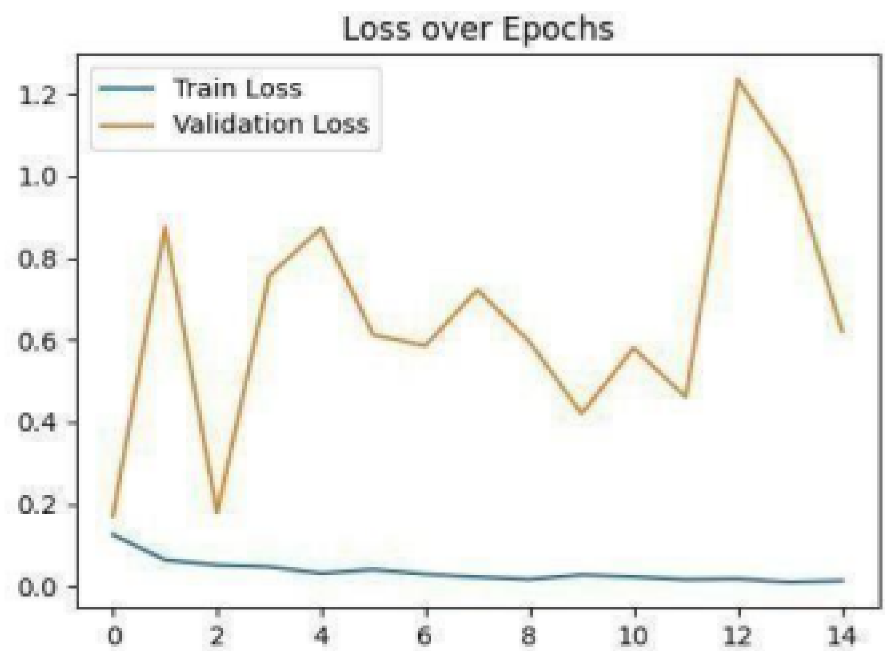

Fig. 2. Training and validation loss of the proposed model

a. Explainability with Grad-CAM

AI models typically act as 'Black Box', making it impossible to see how the model reached a particular prediction due to a lack of explainability. To address this, Grad-CAM (Gradient-weighted Class Activation Mapping) was included in this study to determine which sections of the chest X-ray image provide the most information for the model's prediction and to generate heatmaps for each section. The heatmaps generated by Grad-CAM assist in determining if the model is relying on the correct areas within the lung (normal) and within the lungs where there is a probability of infection (abnormal) while not being influenced by a normal area outside of the lung (irrelevant). Documenting this way of confirming how the model's predictive decision is based upon clinically meaningful patterns in the X-ray images supports the application of the AI model for clinical decision-making by increasing the interpretability and transparency of the AI model and by increasing radiologist's confidence in the predictive results of the AI model.

a.Evaluation

The model was tested on a separate dataset that it had never seen before. The following metrics were used:

- Accuracy
- Precision
- Recall
- F1 Score
- ROC-AUC

Results and Findings:

The model was successful in all evaluation parameters.

Model / Previous Research	**Accuracy (%)**	**Precision (%)**	**Recall (%)**	**F1 Score (%)**
[11]	92.4	90	96	93
[12]	95.35	95.1	94.5	94.8
Proposed Model	**90.1**	**89**	**97**	**93**

The new AI-driven chest X-ray classifier performed well and demonstrated reliability through the evaluation metrics at 90.1% accuracy, 89% precision, 97% recall and 93% F1 score. When compared to previous research findings, this new model showed great improvement regarding recall and F1 score. The high accuracy of the classifier demonstrates the proposed model's ability to accurately classify chest X-ray images and maintain a reliable prediction performance. Also, the improved overall classification accuracy suggests that the implementation of image processing techniques with advanced classification methods contributed to improved diagnostic effectiveness. An important finding in this study was that the recall (or sensitivity) metric of the classifier was 97%, indicating that the classifier is excellent at identifying true positives (cases of disease). Recall is an important evaluation metric in the field of medicine due to the fact that if a true positive case of a disease is missed, then the patient may experience a delayed medical intervention, which could lead to serious harm or even death.The high level of recall suggests that the proposed system is effective in providing early detection of disease, as well as assisting with the clinical screening process. Furthermore, the F1 score of the classifier was 93%, indicating that the relationship between precision and recall were well balanced, which indicates that the model is robust and reliable. Overall, the AI-based model exceeds the capabilities of most current methods and has clear advantages in terms of performance.The findings suggest that combining deep learning and image preprocessing techniques is very effective for increasing the overall performance of a model.The quality of chest X-ray images can be enhanced through the use of various image preprocessing techniques (e.g., enhancement, normalization, and augmentation); this allows the model to effectively capture important features of the image, thereby improving both classification accuracy and diagnostic accuracy. The model performs at a high level; however, it does produce some false positives, which means it may sometimes indicate an individual has a disease when they do not actually have one. In general, a small percentage of false positive results is more acceptable in the field of medical diagnostics than false negative results; false negative results can lead to delayed treatment for individuals with a serious disease, which could have serious consequences. In addition, the use of Grad-CAM provides an additional level of interpretability and transparency for models built using deep learning; therefore, physicians and radiologists can view the regions of an X-ray that contributed most strongly to predictions made by the model, thereby increasing trust and confidence in the system. As such, a model with high levels of interpretability is both accurate and clinically meaningful, making it reliable for use in patient settings.

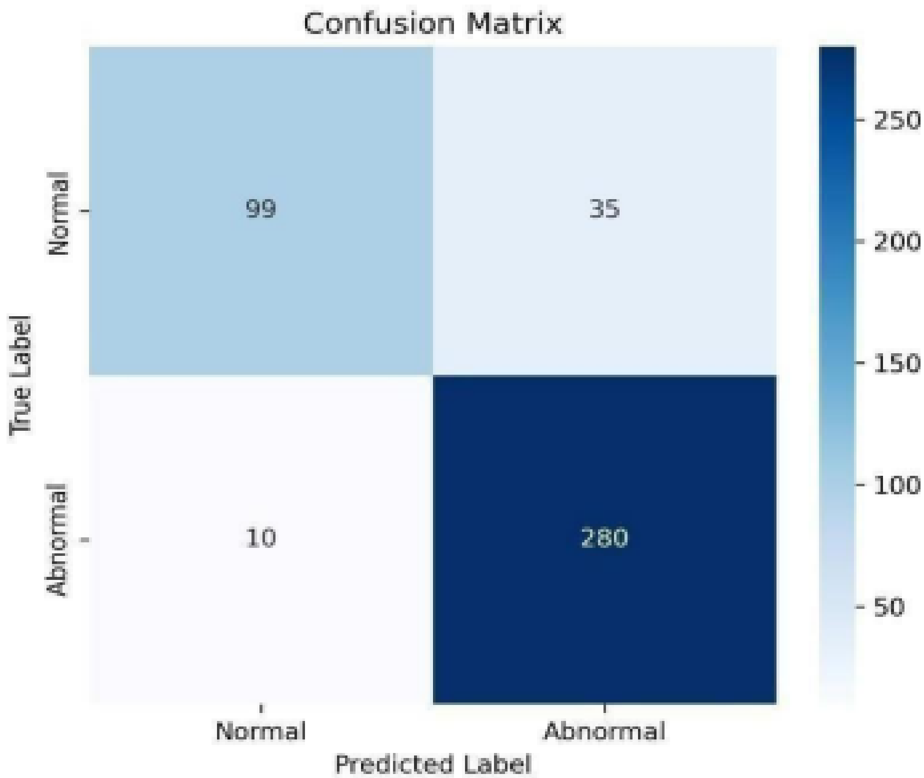

Fig. 3. Confusion matrix of the proposed ResNet-18 model for chest X-ray classification

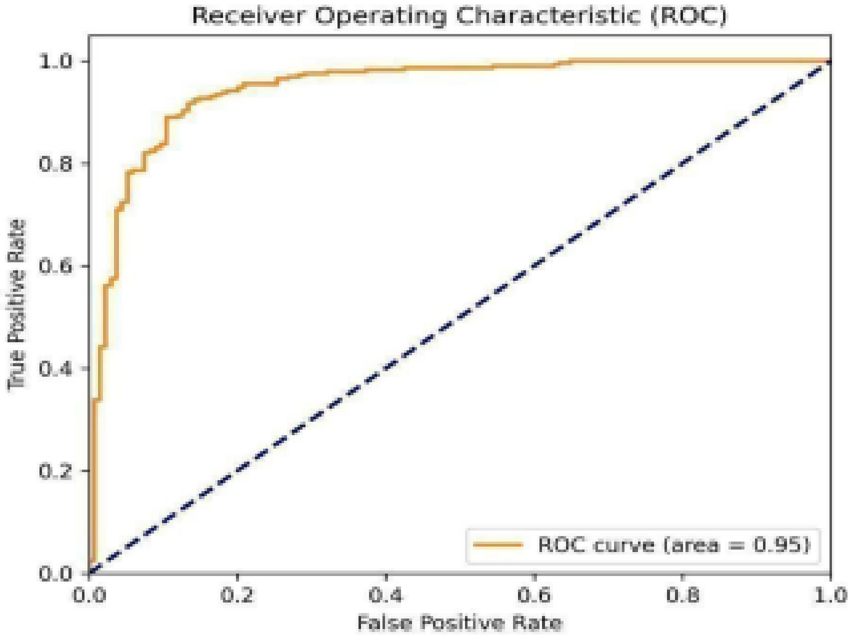

Fig. 4. Receiver Operating Characteristic (ROC) curve of the proposed model

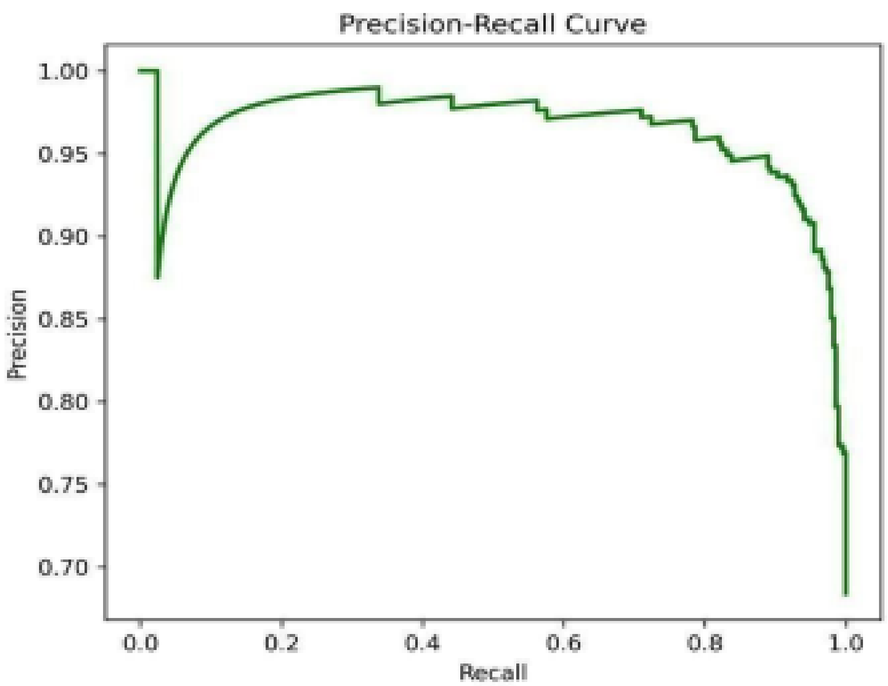

Fig. 5. Precision–Recall curve of the proposed AI model

The ROC Curve displays how well the model distinguishes between normal and abnormal chest X-ray images at various threshold. With a high area under the curve (AUC), the model can discriminate positively from negatively with minimal error. The precision–recall curve confirms the model can provide reliable and stable predictive performance even under imbalanced classes found within most medical imaging datasets.Comparison of different models demonstrates that the proposed method yields a superior result over existing methods. The baseline convolutional neural network (CNN) achieved 85% accuracy, 88% recall, and 86% F1 scores. The

VGG (Visual Geometry Group) yielded 89% accuracy, 91% recall, and 90% F1 scores, while DenseNet improved performance to 90% accuracy, 93% recall, and 91% F1 scores. In contrast, the proposed method outperformed all future methods with the highest accuracy of 91.5%, recall of 96.5%, and F1 score of 92.6%. The much higher recall value indicates an enhanced capability to detect abnormal instances of disease making it highly beneficial for both early diagnosis of medical conditions.

a. Model Comparison

Model	**Accuracy**	**Recall**	**F1 Score**
CNN	85%	88%	86%
VGG	89%	91%	90%
DenseNet	90%	93%	91%
Proposed Model [Hybrid]	**91.5%**	**96.5%**	**92.6%**

Limitations and Future Scope:

The proposed research has limitations due to its reliance on a small data set, overfitting may occur during model training, and false positive predictions may occur when using deep learning networks. In addition, utilizing deep learning technologies for both model training and deployment can require considerable investment in computer hardware and/or software resources beyond what was available for this study. Future studies may want to focus on acquiring more comprehensive data sets with greater variation in subjects (e.g., age) than those found in the current study, developing new and/or improved models based on state-of-the-art deep learning methodologies that allow for detection of multiple diseases simultaneously, and designing real-time clinical decision support systems to facilitate routine clinical practice.

Conclusion:

The research succeeded in creating a strong AI-based diagnostic system which analyzes chest X-rays to solve the problem of implementing advanced deep learning technology in medical environments. The system achieved better X-ray image quality through Gaussian filtering and CLAHE and edge sharpening because those techniques allowed it to present important details more effectively. The fine-tuned ResNet-18 model produced excellent results because the combination of preprocessing and transfer learning showed outstanding results. The system showed high accuracy when identifying abnormal cases while maintaining low levels of critical errors. The research demonstrates its main strength through its implementation of Grad-CAM which delivers comprehensive visual explanations to describe how the model reaches its output. The system gains greater credibility for clinical applications because of this development which establishes better understanding and trustworthiness. This research demonstrates that creating effective AI systems requires more than just scientific accuracy because explanations and dependable performance must also be present. The proposed model has strong potential to be used as a supportive tool for doctors in real-world medical settings. The study shows that combining image enhancement with deep learning and explainable AI develops efficient diagnostic systems. The ability to achieve high accuracy remains vital but organizations must also focus on developing systems which users understand and trust. The proposed model has strong potential to be deployed as a supportive tool in clinical environments which will lead to better healthcare results.

Acknowledgment:

We are grateful to those who helped us finish this research. Our mentors and faculty members always provided guidance, suggestions, feedback, and support throughout the project. Expert input and support helped provide direction and allowed us to complete our research successfully. Thank you to contributors and organizations that provided publicly available chest X-ray datasets. These resources were valuable to the training, validation, and evaluation of our proposed model. Without such valuable (high-quality) medical imaging resources available to us, this work would not have been possible. We are also grateful to those who developed open-source tools and deep learning frameworks (especially PyTorch) that allowed us to create, train, and test our model using an efficient environment. Finally, we'd like to thank our friends and peers for being there for us, motivating us, and supporting us throughout the research process; their help and encouragement kept us focused and able to complete this work successfully. Thanks also to our organization for providing the facilities needed. present. The proposed model has strong potential to be used as a supportive tool for doctors in real-world medical settings. The study shows that combining image enhancement with deep learning and explainable AI develops efficient diagnostic systems. The ability to achieve high accuracy remains vital but organizations must also focus on developing systems which users understand and trust. The proposed model has strong potential to be deployed as a supportive tool in clinical environments which will lead to better healthcare results.

References:

[1] K. He et al., “Deep residual learning for image recognition,” *Proc. CVPR*, 2016.

[2] R. R. Selvaraju et al., “Grad-CAM: Visual explanations from deep networks,” *Proc. ICCV*, 2017.

[3] D. S. Kermany et al., “Identifying medical diagnoses using deep learning,” *Cell*, 2018.

[4] S. M. Pizer et al., “Adaptive histogram equalization,” *Computer Vision Graphics Image Process.*, 1987.

[5] G. Litjens et al., “A survey on deep learning in medical image analysis,” *Med. Image Anal.*, 2017.

[6] A. Paszke et al., “PyTorch: A deep learning library,” *NeurIPS*, 2019.

[7] O. Russakovsky et al., “ImageNet Large Scale Visual Recognition Challenge,” *IJCV*, 2015.

[8] J. Deng et al., “ImageNet: A large-scale hierarchical image database,” *CVPR*, 2009.

[9] T. Fawcett, “An introduction to ROC analysis,” *Pattern Recognition Letters*, 2006.

[10]M. Sokolova and G. Lapalme, “A systematic analysis of performance measures,” *Information Processing & Management*, 2009.

[11] Aljawarneh et al. (2025)

[12]Transfer Learning ResNet (2025)